

assignment1

October 10, 2023

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
from sklearn.preprocessing import MinMaxScaler

data = pd.read_csv('Melbourne_housing.csv')
```

C:\Users\alexm\AppData\Local\Temp\ipykernel_24224\631711095.py:8: DtypeWarning: Columns (13) have mixed types. Specify dtype option on import or set low_memory=False.

```
data = pd.read_csv('Melbourne_housing.csv')
```

Adding imports for graphing and data exploration. Loading in the data using pandas

1 Data Loading and Initial Exploration (20 points):

- 1.1 • Provide information on the dataset, including the number of rows and columns.

```
[2]: data.shape
data.describe
```

```
[2]: <bound method NDFrame.describe of
Rooms Type Method SellerG \ Suburb Address
0 Abbotsford 68 Studley St 2 h SS Jellis
1 Airport West 154 Halsey Rd 3 t PI Nelson
2 Albert Park 105 Kerferd Rd 2 h S hockingstuart
3 Albert Park 85 Richardson St 2 h S Thomson
4 Alphington 30 Austin St 3 h SN McGrath
...
34852 Reservoir 18 Elinda Pl 3 u SP RW
34853 Roxburgh Park 14 Stainsby Cr 4 h S Raine
34854 Springvale South 8 Bellbird Ct 4 h PI Barry
34855 Springvale South 30 Waddington Cr 3 h S Harcourts
34856 Westmeadows 42 Pascoe St 4 h S Barry
```

	Date	Distance	Postcode	Bedroom	...	Landsize	BuildingArea	\
0	3/9/2016	2.5	3067.0	2.0	...	126.0	inf	
1	3/9/2016	13.5	3042.0	3.0	...	303.0	225	
2	3/9/2016	3.3	3206.0	2.0	...	120.0	82	
3	3/9/2016	3.3	3206.0	2.0	...	159.0	inf	
4	3/9/2016	6.4	3078.0	3.0	...	174.0	122	
...
34852	30/09/2017	12.0	3073.0	3.0	...	NaN	105.0	
34853	30/09/2017	20.6	3064.0	4.0	...	NaN	225.0	
34854	30/09/2017	22.2	3172.0	4.0	...	534.0	152.0	
34855	30/09/2017	22.2	3172.0	3.0	...	544.0	NaN	
34856	30/09/2017	16.5	3049.0	4.0	...	813.0	140.0	

	YearBuilt		CouncilArea	Latitude	Longitude	\
0	NaN		Yarra City Council	-37.80140	144.99580	
1	2016.0		Moonee Valley City Council	-37.71800	144.87800	
2	1900.0		Port Phillip City Council	-37.84590	144.95550	
3	NaN		Port Phillip City Council	-37.84500	144.95380	
4	2003.0		Darebin City Council	-37.78180	145.01980	
...
34852	1990.0		Darebin City Council	-37.69769	145.02332	
34853	1995.0		Hume City Council	-37.63665	144.92976	
34854	1970.0	Greater Dandenong City Council		-37.97037	145.15449	
34855	NaN	Greater Dandenong City Council		-37.97751	145.14813	
34856	1960.0		Hume City Council	-37.67631	144.89409	


	Regionname	Propertycount	ParkingArea	Price
0	Northern Metropolitan	4019.0	Carport	NaN
1	Western Metropolitan	3464.0	Detached Garage	840000.0
2	Southern Metropolitan	3280.0	Attached Garage	1275000.0
3	Southern Metropolitan	3280.0	Indoor	1455000.0
4	Northern Metropolitan	2211.0	Parkade	NaN
...
34852	Northern Metropolitan	21650.0	Parkade	475000.0
34853	Northern Metropolitan	5833.0	Underground	591000.0
34854	South-Eastern Metropolitan	4054.0	Carport	NaN
34855	South-Eastern Metropolitan	4054.0	Detached Garage	780500.0
34856	Northern Metropolitan	2474.0	Attached Garage	791000.0

[34857 rows x 22 columns]>

Answer: 22 Columns, 34,857 rows

1.2 • Briefly describe the target variable (e.g., 'Price') and its distribution.

```
[3]: data.Price.describe()  
data.Price.hist(bins=50, figsize=(30,20))  
plt.show()
```



Finding information about the Price variable. Here we can see the count, the mean and all the percentiles and the standard deviation. As we can see on the graph, it is left skewed.

1.3 • Display summary statistics and data types of the features.

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 34857 entries, 0 to 34856  
Data columns (total 22 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   Suburb      34857 non-null  object  
1   Address     34857 non-null  object  
2   Rooms       34857 non-null  int64  
3   Type        34857 non-null  object  
4   Method      34857 non-null  object  
5   SellerG     34857 non-null  object
```

```

6   Date          34857 non-null object
7   Distance      34856 non-null float64
8   Postcode      34856 non-null float64
9   Bedroom       26640 non-null float64
10  Bathroom      26631 non-null float64
11  Car           26129 non-null float64
12  Landsize      23047 non-null float64
13  BuildingArea  13760 non-null object
14  YearBuilt     15551 non-null float64
15  CouncilArea   34854 non-null object
16  Latitude      26881 non-null float64
17  Longitude     26881 non-null float64
18  Regionname    34857 non-null object
19  Propertycount 34854 non-null float64
20  ParkingArea   34857 non-null object
21  Price         27247 non-null float64
dtypes: float64(11), int64(1), object(10)
memory usage: 5.9+ MB

```

1.4 • Identify any missing values and outline a plan to handle them.

```

[5]: #Show sums of NA/NaN
data.isna().sum()

```

```

[5]: Suburb          0
Address            0
Rooms             0
Type              0
Method            0
SellerG           0
Date              0
Distance          1
Postcode          1
Bedroom          8217
Bathroom         8226
Car              8728
Landsize         11810
BuildingArea     21097
YearBuilt        19306
CouncilArea       3
Latitude         7976
Longitude        7976
Regionname        0
Propertycount     3
ParkingArea       0
Price            7610
dtype: int64


```


1.4.1 One way to handle missing data would be Listwise and pairwise deletion techniques. These techniques discard cases during an analysis if they contain missing data. - Taken from “Out of Sight not out of Mind”

2 Exploratory Data Analysis (EDA) (30 points):

2.1 • Visualize the distribution of numeric variables using histograms and box plots.

```
[6]: data.hist(bins=50, figsize=(30,20))  
data.plot(kind='box', figsize=(30,20))  
plt.show()
```
















2.2 • Explore relationships between features and the target variable using scatter plots and correlation matrices.

```
[7]: variables = ['Rooms', 'Distance', 'Postcode', 'Bedroom', 'Bathroom', 'Car', 'Landsize', 'YearBuilt', 'Latitude', 'Longitude', 'Propertycount']

# Create a figure with subplots
fig, axes = plt.subplots(len(variables), 1, figsize=(10, 6*len(variables)))


for i, var in enumerate(variables):
    axes[i].scatter(x=var, y='Price', data=data, edgecolor='b')
    axes[i].set_xlabel(var)
    axes[i].set_ylabel('Price')
    axes[i].set_title(f'{var} v Price')

plt.show()
```



```
[8]: variables = ['Suburb', 'Date', 'Address', 'Type', 'Method', 'SellerG', '
↳ 'BuildingArea', 'CouncilArea', 'Regionname', 'ParkingArea']
for var in variables:
    data[var] = data[var].astype('category')
plt.figure(figsize=(16, 10))
sns.heatmap(data.select_dtypes(exclude=['category']).corr(), annot=True)
plt.title('Correlation')
```

```
[8]: Text(0.5, 1.0, 'Correlation')
```



In Conclusion, the Variable that has the highest effect on price is the amount of Rooms, followed by the Bedroom and Bathroom Count.

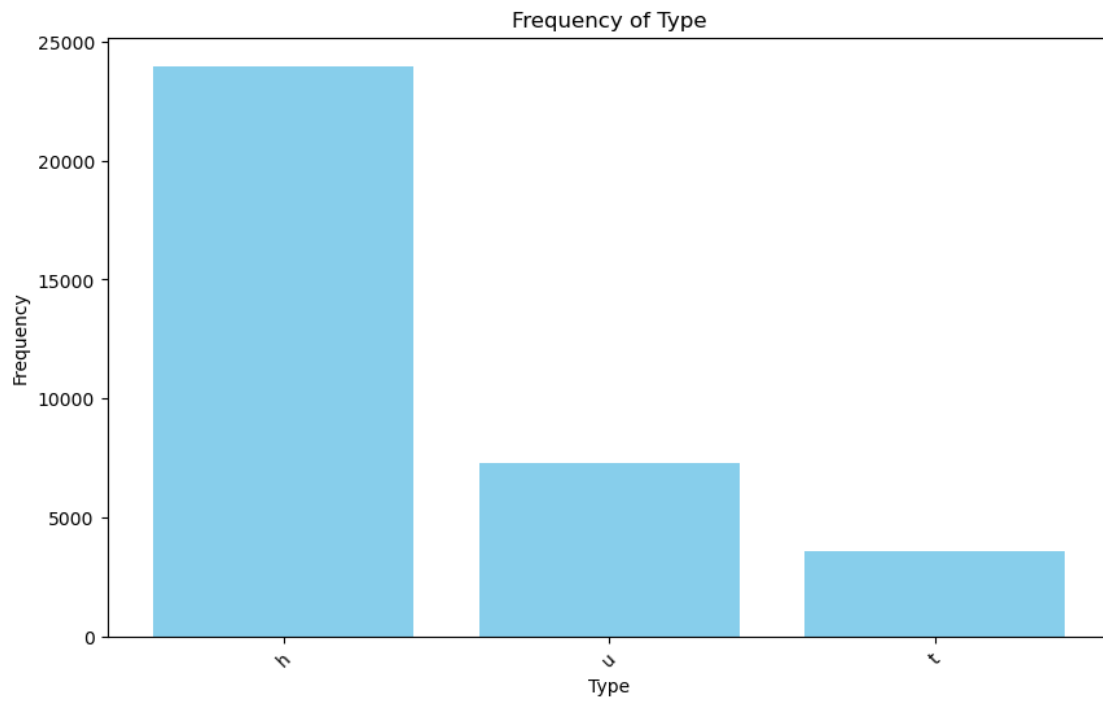
2.3 • Examine categorical variables with bar plots and frequency tables.

```
[9]: categorical_variables = ['Suburb', 'Type', 'Method', 'CouncilArea', '
↳ 'Regionname']

# Loop through each categorical variable
```

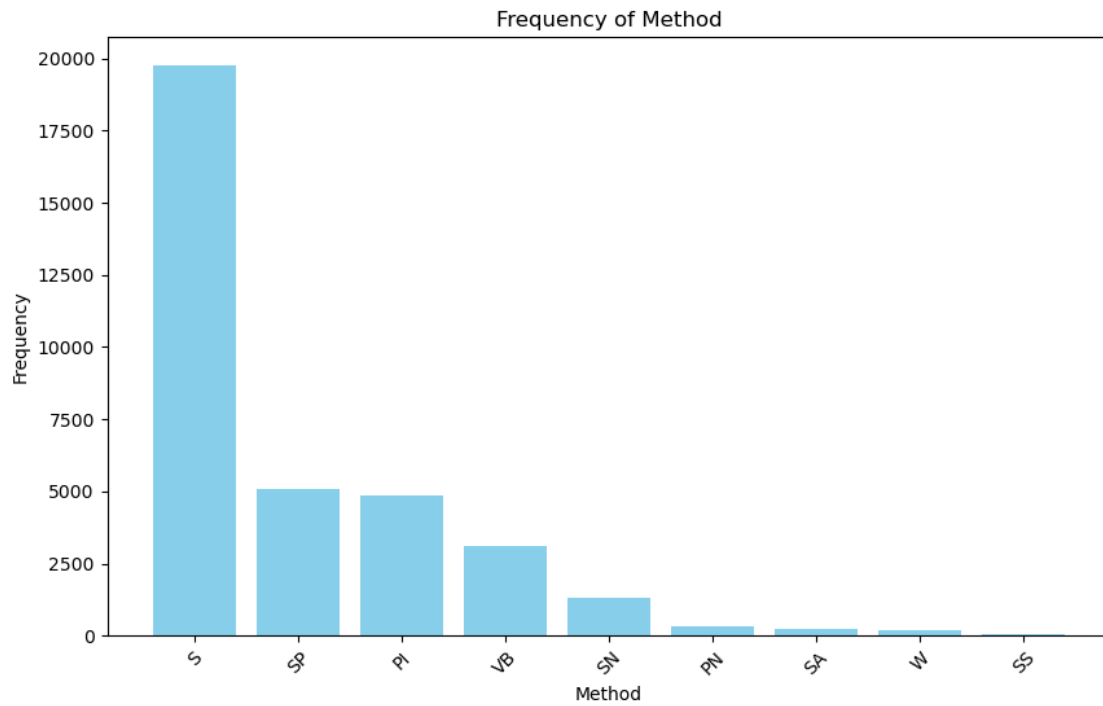

Preston	485
...	
Kalkallo	1
Menzies Creek	1
Monbulk	1
Olinda	1
viewbank	1

Name: count, Length: 351, dtype: int64

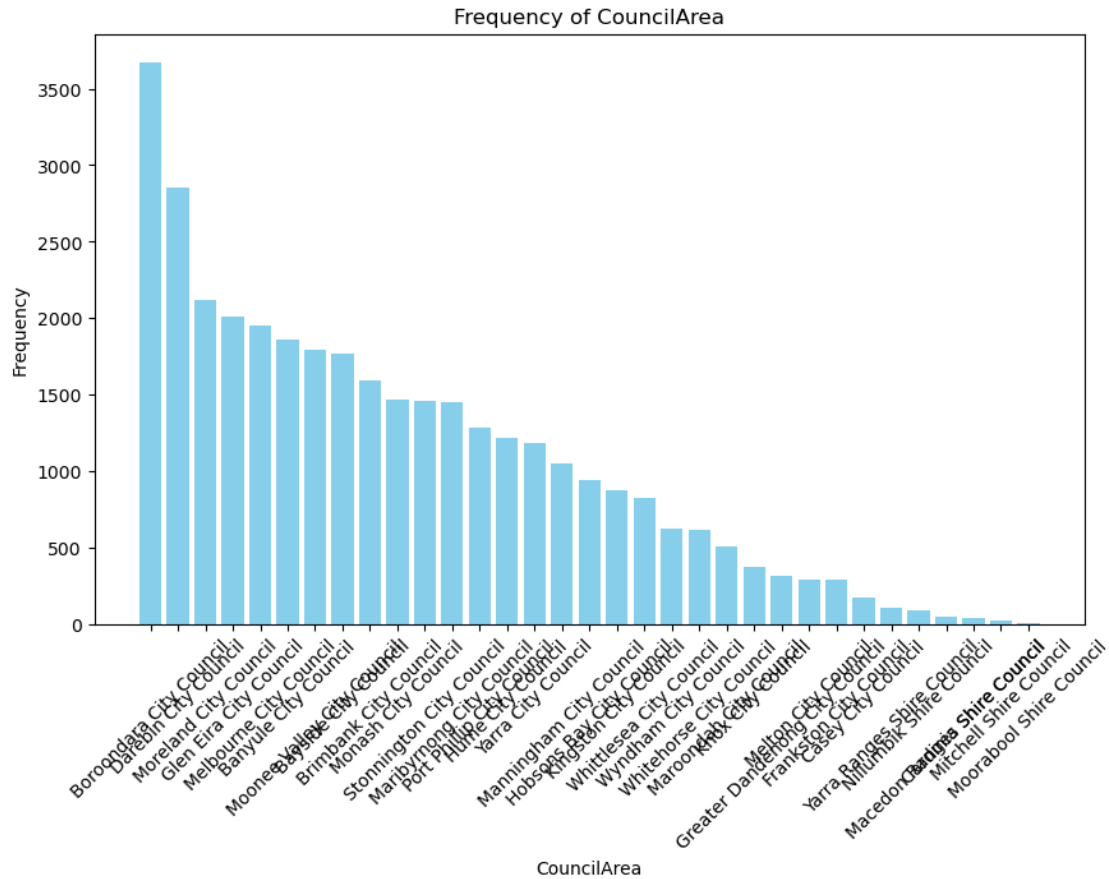


Type	
h	23980
u	7297
t	3580

Name: count, dtype: int64




```
Method
S      19744
SP      5095
PI      4850
VB      3108
SN      1317
PN       308
SA       226
W        173
SS         36
Name: count, dtype: int64
```



CouncilArea	
Boroondara City Council	3675
Darebin City Council	2851
Moreland City Council	2122
Glen Eira City Council	2006
Melbourne City Council	1952
Banyule City Council	1861
Moonee Valley City Council	1791
Bayside City Council	1764
Brimbank City Council	1593
Monash City Council	1466
Stonnington City Council	1460
Maribyrnong City Council	1451
Port Phillip City Council	1280
Hume City Council	1214
Yarra City Council	1186
Manningham City Council	1046
Hobsons Bay City Council	942
Kingston City Council	871
Whittlesea City Council	828

Wyndham City Council	624
Whitehorse City Council	618
Maroondah City Council	506
Knox City Council	371
Greater Dandenong City Council	314
Melton City Council	292
Frankston City Council	290
Casey City Council	176
Yarra Ranges Shire Council	102
Nillumbik Shire Council	88
Macedon Ranges Shire Council	46
Cardinia Shire Council	41
Mitchell Shire Council	20
Moorabool Shire Council	7

Name: count, dtype: int64



Regionname	
Southern Metropolitan	11836
Northern Metropolitan	9560
Western Metropolitan	6799

Eastern Metropolitan	4377
South-Eastern Metropolitan	1739
Eastern Victoria	228
Northern Victoria	203
Western Victoria	115

Name: count, dtype: int64

Here we can see the most frequent categorical Variables

2.4 • Identify potential outliers and discuss their impact on the dataset.

Some features that have outliers are: Price and Landsize, these can be seen in the BoxPlots above as the largest plots and are clearly affected by outliers.

Outliers can significantly affect measures like mean and standard deviation. The mean is particularly sensitive to outliers. Outliers can also lead to skewed distributions, which may not accurately represent the underlying data. Outliers may also indicate data quality issues, such as errors in data entry.

3 Feature Engineering (40 points):

3.1 • Apply at least five feature engineering techniques to improve the dataset for modeling purposes. Some ideas include:

```
[10]: num_cols = data.select_dtypes(include=["number"]).columns
      cat_cols = data.select_dtypes(include=["category"]).columns

      #To handle missing numerical data, we can use Mean/Median Imputation:
      for num in num_cols:
          data[num].fillna(data[num].mean(), inplace=True)

      # Encoding categorical variables with one-hot encoding
      def encode_and_bind(data: pd.DataFrame, feature: str) -> None:
          dummies = pd.get_dummies(data[[feature]])
          res = pd.concat([data, dummies], axis=1)
          res = res.drop([feature], axis=1)
          return res

      for cat in cat_cols:
          data = encode_and_bind(data, cat)

      # Removing extreme values
      for col in num_cols:
          z_scores = np.abs(stats.zscore(data[col]))
          outliers = data[(z_scores > 3)]
          data = data.drop(outliers.index)

      # Scaling numeric features using normalization
```

```
scaler = MinMaxScaler()

mask = (data == np.inf) | (data == -np.inf)
data = data[~mask.any(axis=1)]

data[num_cols] = scaler.fit_transform(data[num_cols])

# Creating an interaction feature
data["PriceLandsize"] = data["Landsize"] * data["Price"]
```

4 Conclusion and Recommendations (10 points):

4.1 • Summarize the key findings from the EDA and feature engineering processes.

4.1.1

EDA: The dataset contains information about properties in Melbourne, including features like location, type of property, number of rooms, land size, building area, etc. The data originally contains a lot of missing or NA values that can often change the outcome of graphs etc. The main feature that has a big effect on Price is the no. of rooms.

FE: It is rather easy to remove numerical values that are missing and fill them in with the medians of the columns. Creating interaction features can be great for visualising effects of one feature on another.

Overall I believe that the Melbourne Housing Dataset is a great dataset to learn on and try out new things on as it features a lot of missing values and outliers that can be great for perfecting EDA and FE processes.