

# assignment2

October 17, 2023

## 1 1. Data Exploration and Preparation (15 points):

- Load the Iris dataset and examine its structure.
- Explore the dataset by calculating summary statistics and visualizing the data using appropriate plots (e.g., histograms, scatter plots).
- Split the dataset into training (80%) and testing sets (20%).

```
[11]: %matplotlib inline
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.tree import plot_tree, DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd

#loading Data
data = load_iris()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target
```

```
[12]: #Exploring dataset
df.shape
df.describe
```

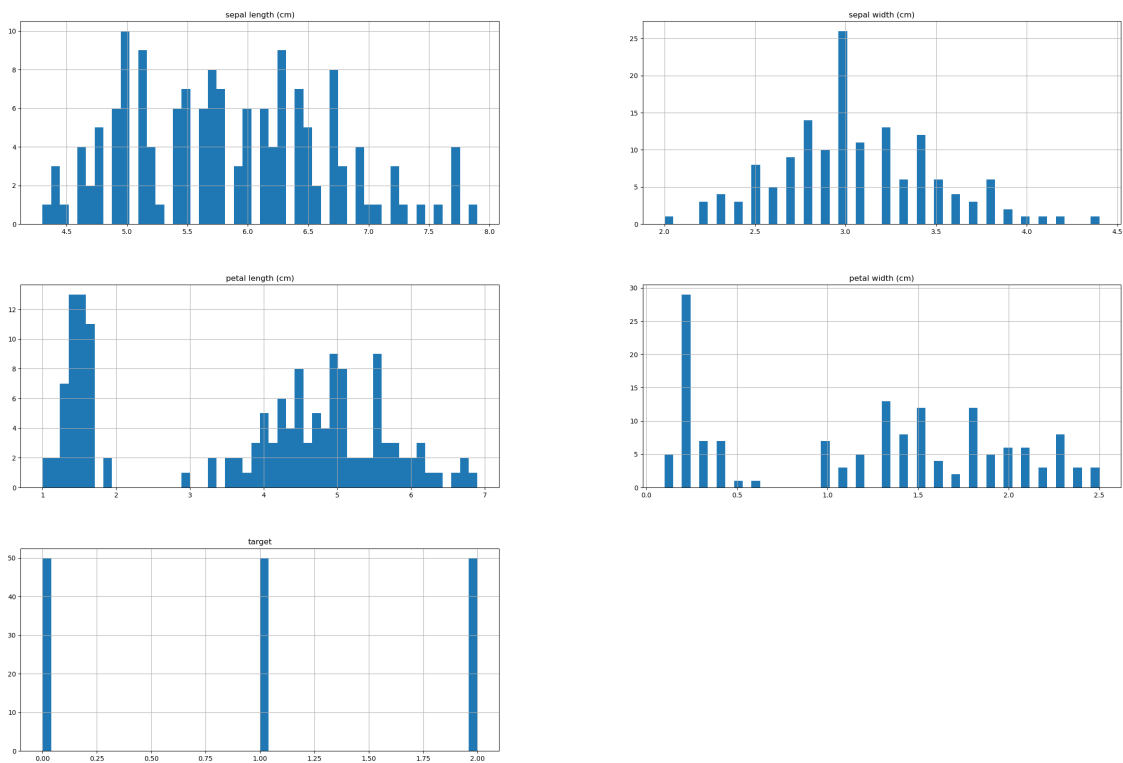
```
[12]: <bound method NDFrame.describe of
petal length (cm)  petal width (cm)  \
0                5.1                3.5                1.4                0.2
1                4.9                3.0                1.4                0.2
2                4.7                3.2                1.3                0.2
3                4.6                3.1                1.5                0.2
4                5.0                3.6                1.4                0.2
..                ...                ...                ...                ...
145               6.7                3.0                5.2                2.3
146               6.3                2.5                5.0                1.9
147               6.5                3.0                5.2                2.0
```

148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

	target
0	0
1	0
2	0
3	0
4	0
..	...
145	2
146	2
147	2
148	2
149	2

[150 rows x 5 columns]>

```
[13]: df.hist(bins=50, figsize=(30,20))
plt.show()
```



```
[14]: plt.figure(figsize=(12, 8))
```

```

# Define a list of colors for the different classes
colors = ['b', 'r', 'g']

# Define the class labels
classes = df['target'].unique()


# Create scatter plots for each pair of features
for i in range(len(classes)):
    class_data = df[df['target'] == classes[i]]
    plt.scatter(class_data['sepal length (cm)'], class_data['sepal width_
↵(cm)'], c=colors[i], label=f'Class {classes[i]}')

# Set labels and title
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Sepal Length vs Sepal Width')

# Add a legend
plt.legend()

# Show the plot
plt.show()

```



```
[15]: #splitting the dataset into testing and training sets
X_train, X_test, Y_train, Y_test = train_test_split(df[data.feature_names],
↳df['target'], test_size=0.2, random_state=0)
```

We can clearly see that the classes have a correlation with the sepal with and length.

We can also see that sepal length and width are left skewed.

## 2 2. Decision Tree Implementation (30 points):

- Implement a decision tree classifier using scikit-learn's DecisionTreeClassifier class.
- Train the decision tree classifier on the training data.
- Set and tune hyperparameters like the maximum depth of the tree, minimum samples per leaf, or other relevant parameters.

```
[16]: #Train the decision tree on the training data
#Hyperparameter Tuning
tree_classifier_tuned = DecisionTreeClassifier(max_depth=3, min_samples_leaf=4,
↳random_state=0)
tree_classifier_tuned.fit(X_train, Y_train)
y_pred_tuned = tree_classifier_tuned.predict(X_test)
accuracy_tuned = accuracy_score(Y_test, y_pred_tuned)
print(f'Tuned Accuracy: {accuracy_tuned:.2f}')

#Predict on the test data
Y_pred = tree_classifier_tuned.predict(X_test)

#Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, Y_pred)
print(f'Default Accuracy: {accuracy:.2f}')
```

Tuned Accuracy: 1.00

Default Accuracy: 1.00

Here we can see the effect of tuning on the accuracy of the Model, in this case the default accuracy was at 100% so there were no improvements to be made, however we can still experiment as seen to tweak the accuracy

## 3 3. Model Evaluation (20 points):

- Use the trained classifier to predict the species of iris flowers in the testing dataset.
- Evaluate the model's performance using appropriate classification metrics such as accuracy, precision, recall, F1-score, etc.

```
[17]: #The training was done in section 2. on line 6 with the following code:
#Y_pred = tree_classifier.predict(X_test)
```

```

#Evaluating the Model Performance
#Re-Evaluate accuracy
accuracy = accuracy_score(Y_test, Y_pred)
print(f'Accuracy: {accuracy:.2f}')

#Evaluate precision
precision = precision_score(Y_test, Y_pred, average='weighted')
print(f'Precision: {precision:.2f}')

#Evaluate recall
recall = recall_score(Y_test, Y_pred, average='weighted')
print(f'Recall: {recall:.2f}')

#Evaluate F1-score
f1 = f1_score(Y_test, Y_pred, average='weighted')
print(f'F1 Score: {f1:.2f}')

#Generate a full classification report
class_report = classification_report(Y_test, Y_pred, target_names=data.
    ↪target_names)
print(f'Classification Report:\n{class_report}')

```

Accuracy: 1.00

Precision: 1.00

Recall: 1.00

F1 Score: 1.00

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	11
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

## 4 4. Visualization (15 points):

- Visualize the decision tree structure that you've built, showing how it makes decisions based on the features.

```

[18]: #Visualize the decision tree

#Converting to list so that plot_tree accepts the set
class_names = data.target_names.tolist()

```

