


# Training a simple neural network



## Table of Contents

### Training a simple neural network

Define a generic NN layer

Define a network as a concatenation of many layers

Create a two-layer network

Train on random data

Visualize Network's Performance

```
1 begin
2   using Random
3   using PlutoUI
4   using Latexify
5   TableOfContents()
6 end
```

```

1 begin
2     using Plots
3
4     # Packages for automatic differentiation and neural networks
5     using Flux, Zygote
6 end

```

## Define a generic NN layer

Use a struct to define a generic layer of a neural network. The struct fields comprise:

- $W$ : a weight matrix of floats connecting the layer's input to its output
- $b$ : a vector of float biases which serve to modulate the default output of each neuron
- `activation`: an activation function that maps the layer's input to its output
- a constructor, which takes the input and output dimensions of the layer as parameters along with an activation function (default is the identity function) and randomly initializes the weights and biases

The final expression in the block below calls the `Layer` struct as a function with an input vector as an argument and returns an output vector, effectively it implements the feedforward step for the layer.

```

1 begin
2     struct Layer
3         W::Matrix{Float32} # weight matrix - Float32 for faster gradients
4         b::Vector{Float32} # bias vector
5         activation::Function
6         Layer(in::Int64, out::Int64, activation::Function=identityFunction) =
7             new(randn(out, in), randn(out), activation) # constructor
8     end
9
10    (m::Layer)(x) = m.activation.(m.W * x .+ m.b) # feed-forward pass
11 end

```

Define some required activation functions. `ReLU` is a standard neural network activation function.

```

1 begin
2     ReLU(x) = max(0, x)
3     identityFunction(x) = x
4 end;

```

## Define a network as a concatenation of many layers

Again, we use a struct to define a network comprising an arbitrary number of layers. We need just one field, `layers`, which is a vector of type `Layer`. The constructor takes a variable number of `Layer` arguments and assigns them to the `layers` field.

The final function definition in the block serves to propagate its vector argument `x` through the entire network. For clarity, the expression `reduce((left,right)->right∘left, m.layers)(x)` can be broken down into a number of elements:

- `right∘left` represents the composition of the function `right` and `left`. So `(right∘left)(x)` is equivalent to `right(left(x))`.
- `(left, right)->right∘left` is an anonymous function taking two arguments: `left` and `right`, which executes a composition of its functional arguments
- `reduce` is a function that applies its first argument, a function, to the elements of a collection. In this case the layers of the network, which is equivalent to `right(left(x))`.

```

1 begin
2     struct Network
3         layers::Vector{Layer}
4         Network(layers::Vararg{Layer}) = new(vcat(layers...))
5         # constructor - allow arbitrarily many layers
6     end
7
8     (n::Network)(x) = reduce((left, right) -> right ∘ left, n.layers)(x)
9     # perform layer-wise operations over arbitrarily many layers
10 end

```

## Create a two-layer network

Define a neural network with two hidden layers of 100 neurons each, and one input and one output unit.


The 1st hidden layer uses a rectified linear unit (ReLU) activation function, and the 2nd uses the default identity function.

```
Network([Layer(100×7 Matrix{Float32}:
```

```
1.22684 -1.22644 -2.72411 -1.65359 -0.790929 1.12893 -
```

```
1 begin
2   inputs = rand([-1, 1], (7, 10)) # 7-bit input, 10 examples
3   targetOutput = rand([-1, 1], (10, 10)) # 10-bit output, 10 examples
4
5   mse(x, y) = sum((x .- y).^2) / length(x) # MSE will be our loss function
6
7   Random.seed!(54321) # for reproducibility
8
9   twoLayerNeuralNet = Network(Layer(7, 100, ReLu), Layer(100, 10))
10   # instantiate a two-layer network
11 end
```

## ReLU activation function



The error (or loss) function is the mean squared difference between the actual output and target output.

$$\text{mse}(x, y) = \frac{\sum (x - y)^2}{\text{length}(x)}$$

## Train on random data

We use the Flux library functions to calculate the relevant gradients for the `Layer` and `Network` structs.

- we first extract the trainable parameters (weights and biases) from the network
- we assign an optimiser, ADAM, which adjusts the rate at which we change these parameters
- we set up vectors to log the training performance
- then we iterate over the training set, adjusting the weights after each iteration with repeated calls to `Zygote.gradient` followed by weight updates.

```

1 begin
2
3     Flux.@functor Layer      # set the Layer-struct as being differentiable
4     Flux.@functor Network    # set the Network-struct as being differentiable
5
6     parameters = Flux.params(twoLayerNeuralNet)
7         # obtain the parameters of the layers (recurses through network)
8
9     optimizer = ADAM(0.05) # from Flux-library
10
11     netOutput = [] # store output for plotting
12     lossCurve = [] # store loss for plotting
13
14     for i in 1:1000
15         # Randomly select one of the 10 vectors for each training iteration
16         random_index = rand(1:10)
17         single_input = inputs[:, random_index]
18         single_output = targetOutput[:, random_index]
19
20         # Calculate the gradients for the network parameters
21         gradients = Zygote.gradient(
22             () -> mse(twoLayerNeuralNet(single_input), single_output),
23             parameters
24         )
25
26         # Update the parameters using the gradients and optimiser settings.
27         Flux.Optimise.update!(optimizer, parameters, gradients)
28
29         # Log the performance for later plotting
30         actualOutput = twoLayerNeuralNet(single_input)
31         push!(netOutput, actualOutput)
32         push!(lossCurve, mse(actualOutput, single_output))
33     end
34 end

```

## Visualize Network's Performance




```


1 @bind plotIndex Slider(1:10:1000, default=1000)
2

```

## Neural Network fit



## Loss curve




```

1 begin
2   #   outputPlot = scatter(1:10, targetOutput,
3   #       title = "Neural Network fit", label = "Data points", legend=:topleft
4   #   )
5
6   #   plot!(outputPlot, 1:10, netOutput[plotIndex],
7   #       label = "Learned function", lw = 4, color = :red
8   #   )
9
10  #   annotate!(3.0, -0.75, text(plotIndex, :blue, :right, 15))
11
12  # lossCurvePlot = plot(lossCurve, title = "Loss curve", legend=:none)
13
14  #   plot(outputPlot)
15
16  # Bar plot for the data points
17  outputPlot = bar(1:10, targetOutput,
18      title = "Neural Network fit", label = "Data points", legend=:topleft
19  )
20
21  # Bar plot for the learned function
22  bar!(outputPlot, 1:10, netOutput[plotIndex],
23      label = "Learned function", lw = 4, color = :red
24  )
25
26  # Annotation
27  annotate!(outputPlot, 3.0, -0.75, text(string(plotIndex), :blue, :right, 15))
28
29  # Line plot for the loss curve
30  lossCurvePlot = plot(lossCurve, title = "Loss curve", legend=:none)
31
32  # Combine plots
33  plot(outputPlot, lossCurvePlot, layout = (2, 1))
34 end

```

# Loss curve



```
1 plot(lossCurvePlot)
```