



## Table of Contents

### Sentiment analysis using a transformer

Model training

Download module from GitHub

Setup the training data

Tokenisers

Tokenise

Model definition

Training

Save the model

# Sentiment analysis using a transformer

---

## Model training

---

### Tip

Hidden below is a useful snippet of HTML to setup a `restart` button in case training gets out of hand.

Restart

## Download module from GitHub

Because we're working from a GitHub repo and not the standard Julia repository, we have to manage the installation and use of all packages rather than rely on Pluto.

```

1 begin
2     import Pkg
3     Pkg.activate(mktempdir())
4     Pkg.develop(url="https://github.com/rgreilly/Transformers")
5     Pkg.add(["Revise", "PlutoUI", "Flux", "DataFrames", "Printf",
6             "BSON", "JSON", "Arrow", "StatsBase", "Unicode", "Random",
7             "DataStructures", "ProgressMeter", "RemoteFiles"])
8
9     using Revise
10    using TransformersLite
11    using PlutoUI
12    using Flux
13    using Flux.CUDA
14    using Flux: DataLoader
15    using DataFrames
16    using BSON, JSON
17    using Arrow
18    using Printf
19    using StatsBase
20    using StatsBase: mean
21    using Dates
22    using Unicode
23    using Random
24    using DataStructures
25    using TransformersLite
26    using RemoteFiles
27 end;

```

```

\ Activating new project at 'C:\Users\alexm\AppData\Local\Temp\jl_BQHFa'
\
\ Cloning git-repo 'https://github.com/rgreilly/Transformers'
Path 'C:\Users\alexm\.julia\dev\TransformersLite' exists and looks like the
correct repo. Using existing path.
Resolving package versions...
Updating 'C:\Users\alexm\AppData\Local\Temp\jl_BQHFa\Project.toml'
[6579f8b0] + TransformersLite v0.1.0 'C:\Users\alexm\.julia\dev\Transforme
rsLite'
Updating 'C:\Users\alexm\AppData\Local\Temp\jl_BQHFa\Manifest.toml'
[621f4979] + AbstractFFTs v1.5.0
[79e6a3ab] + Adapt v3.7.2
[dce04be8] + ArgCheck v2.3.0
[69666777] + Arrow v2.7.0
[31f734f8] + ArrowTypes v2.3.0
[a9b6321e] + Atomix v0.1.0
[ab4f0b2a] + BFloat16s v0.4.2
[fbb218c0] + BSON v0.3.7
[198e06fe] + BangBang v0.3.39
[9718e550] + Baselet v0.1.1
[c3b6d118] + BitIntegers v0.3.1
[fa961155] + CEnum v0.4.2
[052768ef] + CUDA v4.4.1
[1af6417a] + CUDA_Runtime_Discovery v0.2.2
[082447d4] + ChainRules v1.58.1
[d360d2e6] + ChainRulesCore v1.19.0
[5ba52731] + CodecLz4 v0.4.1
[6b39b394] + CodecZstd v0.8.1
[bbf7d656] + CommonSubexpressions v0.3.0
[34da2185] + Compat v4.10.1
[a33af91c] + CompositionsBase v0.1.2
[f0e56b4a] + ConcurrentUtilities v2.3.0
[187b0558] + ConstructionBase v1.5.4
[6add18c4] + ContextVariables v0.1.3

```

In addition to the list of modules, we also need to include individual Julia files from the repo. This is done using the `RemoteFiles` module. However, this downloads them as JSON objects, which we need to convert back to regular `.jl` files.

```

1 begin
2     @RemoteFileSet FILES "Transformer utilities" begin
3
4         utilities = @RemoteFile
5             "https://github.com/rgreilly/Transformers/blob/main/examples/utilities.jl"
6             dir="utilities" file="utilities.jl.json"
7
8         training = @RemoteFile
9             "https://github.com/rgreilly/Transformers/blob/main/examples/training.jl"
10            dir="utilities" file="training.jl.json"
11
12     end
13
14     download(FILES) # Files downloaded in JSON format
15 end

```

`convertJSON` (generic function with 1 method)

```

1 function convertJSON(inFile, outFile)
2     body = JSON.parsefile(inFile)["payload"]["blob"]["rawLines"]
3     open(outFile, "w") do f
4         for i in body
5             println(f, i)
6         end
7     end
8 end

```

```

1 begin
2     convertJSON("utilities/utilities.jl.json", "utilities/utilities.jl")
3     convertJSON("utilities/training.jl.json", "utilities/training.jl")
4
5     include("utilities/utilities.jl")
6     include("utilities/training.jl")
7 end;

```

## Setup the training data

- Setup the file path to the Kaggle Amazon reviews dataset
- Assign values to various hyper-parameters and store them in a dictionary.
- Set number of training epochs

### Tip

Here is where you can manipulate various training parameters - `pdrop`: proportion of weights to dropout (i.e., set to zero); `dim-embedding`: size of embedding; `n_epoch`: number of epochs.

```

1 begin
2   path = normpath(joinpath(@__DIR__, "..", "assignment6/datasets",
3     "amazon_reviews_multi", "en", "1.0.0"))
4   filename = "train.arrow"
5   to_device = cpu # gpu or cpu
6
7   filepath = joinpath(path, filename)
8
9   df = DataFrame(Arrow.Table(filepath))
10  display(first(df, 20))
11  println("")
12
13  hyperparameters = Dict(
14    "seed" => 314159,
15    "tokenizer" => "none", # options: none bpe affixes
16    "nlabels" => 5,
17    "pdrop" => 0.75, #CHANGE HERE FROM 0.1->0.5->0.75
18    "dim_embedding" => 128 #CHANGE HERE FROM 32->64->128
19  )
20  nlabels = hyperparameters["nlabels"]
21  n_epochs = 10
22 end;

```

20x9 DataFrame

Row	Column1	review_id	product_id	reviewer_id	star ...
	Int64	String	String	String	Int6 ...
1	200000	en_0964290	product_en_0740675	reviewer_en_0342986	...
2	200001	en_0690095	product_en_0440378	reviewer_en_0133349	...
3	200002	en_0311558	product_en_0399702	reviewer_en_0152034	...
4	200003	en_0044972	product_en_0444063	reviewer_en_0656967	...
5	200004	en_0784379	product_en_0139353	reviewer_en_0757638	...
⋮	⋮	⋮	⋮	⋮	⋮
17	200016	en_0619473	product_en_0250211	reviewer_en_0056679	...
18	200017	en_0533035	product_en_0566399	reviewer_en_0488191	...
19	200018	en_0832890	product_en_0304984	reviewer_en_0667005	...
20	200019	en_0550306	product_en_0387159	reviewer_en_0627216	...

5 columns and 11 rows omitted

## Tokenisers

Select a tokeniser. In this case, none, which just uses the various inflected word forms.

```
1 begin
2   if hyperparameters["tokenizer"] == "bpe"
3     directory = joinpath("vocab", "bpe")
4     path_rules = joinpath(directory, "amazon_reviews_train_en_rules.txt")
5     path_vocab = joinpath(directory, "amazon_reviews_train_en_vocab.txt")
6     tokenizer = load_bpe(path_rules, startsym=".")
7   elseif hyperparameters["tokenizer"] == "affixes"
8     directory = joinpath("vocab", "affixes")
9     path_vocab = joinpath(directory, "amazon_reviews_train_en_vocab.txt")
10    tokenizer = load_affix_tokenizer(path_vocab)
11  elseif hyperparameters["tokenizer"] == "none"
12    path_vocab = joinpath("vocab", "amazon_reviews_train_en.txt")
13    tokenizer = identity
14  end
15
16  vocab = load_vocab(joinpath(@__DIR__, path_vocab))
17  indexer = IndexTokenizer(vocab, "[UNK]")
18
19  display(tokenizer)
20  println("")
21  display(indexer)
22  println("")
23
24 end
```

identity (generic function with 1 method)



IndexTokenizer{String}(length(vocabulary)=6654, unksym=[UNK])

## Tokenise

Extract the review body and star rating from the dataframe and create embeddings. Partition data into training and validation sets.

```

1 begin
2   documents = df[!, :review_body]
3   labels = df[!, :stars]
4   max_length = 50
5   indices_path = joinpath(@__DIR__, "outputs", "indices_" *
6     hyperparameters["tokenizer"] * ".bson")
7   @time tokens = map(d->preprocess(d, tokenizer, max_length=max_length),
8     documents)
9   @time indices = indexer(tokens)
10
11  y_labels = Int.(labels)
12  if nlabels == 1
13    y_labels[labels .≤ 2] .= 0
14    y_labels[labels .≥ 4] .= 1
15    idxs = labels .!= 3
16    y_labels = reshape(y_labels, 1, :)
17  else
18    idxs = Base.OneTo(length(labels))
19    y_labels = Flux.onehotbatch(y_labels, 1:nlabels)
20  end
21
22  X_train, y_train = indices[:, idxs], y_labels[:, idxs];
23  rng = MersenneTwister(hyperparameters["seed"])
24  train_data, val_data = split_validation(X_train, y_train; rng=rng)
25
26  println("train samples:      ", size(train_data[1]), " ", size(train_data[2]))
27  println("validation samples: ", size(val_data[1]), " ", size(val_data[2]))
28  println("")
29 end

```

```

4.023848 seconds (28.85 M allocations: 1.836 GiB, 17.55% gc time, 3.22% compilation time)
12.477690 seconds (31.81 k allocations: 81.916 MiB, 0.33% gc time, 0.53% compilation time)
train samples:      (50, 184500) (5, 184500)
validation samples: (50, 20500) (5, 20500)

```

## Model definition

Assemble the model's components.

### Tip

Here's where you might want to adjust the number and nature of the encoder blocks (e.g., attention heads, dropout), number of Dense layers and their characteristics (e.g., activation function, dimensions), the number of dropout layers.

**MethodError: no method matching TransformersLite.TransformerEncoderBlock(::Int64, ::Int64, ::Int64; nheads::Int64, pdrop::Float64)**

Closest candidates are:

TransformersLite.TransformerEncoderBlock(::Int64, ::Int64, ::Int64; pdrop, act) got unsupported keyword argument "nheads"

@ TransformersLite C:\Users\alexm\.julia\dev\TransformersLite\src\encoder.jl:21

1. `kwerr(::NamedTuple{(:nheads, :pdrop)}, Tuple{Int64, Float64}), ::Type, ::Int64, ::Int64, ::Int64) @ error.jl:165`
2. `top-level scope @ Local: 4`

```

1 begin
2     dim_embedding = hyperparameters["dim_embedding"]
3     pdrop = hyperparameters["pdrop"]
4     model = TransformersLite.TransformerClassifier(
5         Embed(dim_embedding, length(indexer)),
6         PositionEncoding(dim_embedding),
7         Dropout(pdrop),
8         TransformerEncoderBlock[
9             TransformerEncoderBlock(4, dim_embedding, dim_embedding * 4;
10                pdrop=pdrop), #Added ,
11
12                TransformerEncoderBlock(4, dim_embedding, dim_embedding * 4;
13                pdrop=pdrop), #Added extra transformer block normal
14
15                TransformerEncoderBlock(4, dim_embedding, dim_embedding * 4,
16                nheads=8; pdrop=pdrop), #Added extra encoder block with attention Head
17            ],
18            Dropout(pdrop), #Adding Dropout Layers
19            Dense(dim_embedding, 1),
20            FlattenLayer(),
21            Dense(max_length, nlabels),
22            Dropout(pdrop), #Adding Dropout Layers
23            Dense(dim_embedding, 256, relu), #Adding dense layers
24            Dropout(pdrop), #Adding Dropout Layers
25            Dense(256, 128, relu), #Adding dense layers
26        )
27     display(model)
28     println("")
29     model = to_device(model)
30
31     hyperparameters["model"] = "$(typeof(model).name.wrapper)"
32     hyperparameters["trainable parameters"] = sum(length, Flux.params(model));
33
34     if nlabels == 1
35         loss(x, y) = Flux.logitbinarycrossentropy(x, y)
36         accuracy(y_hat, y) = mean((Flux.sigmoid.(y_hat) .> 0.5) .== y)
37     else
38         loss(x, y) = Flux.logitcrossentropy(x, y)
39         accuracy(y_hat, y) = mean(Flux.onecold(y_hat) .== Flux.onecold(y))
40     end
41 end;

```



# Training

- Setup the dataloaders to batch and shuffle the training and validation data.
- Print out initial accuracy and loss values for the validation data.
- Setup a sub-directory in the outputs directory, based on date and time, to store the trained model and associated hyperparameters.
- call the `train!` method and log training progress.

UndefVarError: `model` not defined

```

1 begin
2   opt_state = Flux.setup(Adam(), model)
3   batch_size = 32
4
5   train_data_loader = DataLoader(train_data |> to_device; batchsize=batch_size,
6     shuffle=true)
7   val_data_loader = DataLoader(val_data |> to_device; batchsize=batch_size,
8     shuffle=false)
9
10  val_acc = batched_metric(model, accuracy, val_data_loader)
11  val_loss = batched_metric(model, loss, val_data_loader)
12
13  @printf "val_acc=%.4f%% ; " val_acc * 100
14  @printf "val_loss=%.4f \n" val_loss
15  println("")
16
17  directory2 = normpath( joinpath(@__DIR__, "..", "outputs",
18    Dates.format(now(), "yyyymmdd_HHMM")))
19  mkpath(directory2)
20
21  hyperparameter_path = joinpath(directory2, "hyperparameters.json")
22  open(hyperparameter_path, "w") do f
23    JSON.print(f, hyperparameters)
24  end
25  println("saved hyperparameters to $(hyperparameter_path).")
26  println("")
27
28  start_time = time_ns()
29  history = train!(
30    loss, model, train_data_loader, opt_state, val_data_loader;
31    num_epochs=n_epochs)
32  end_time = time_ns() - start_time
33
34  println("done training")
35  @printf "time taken: %.2fs\n" end_time/1e9
36 end

```

UndefVarError: `accuracy` not defined

```

1 accuracy

```

# Save the model

---

Save model, embeddings, and training history to the outputs sub-directory.

**UndefVarError: `model` not defined**

```
1 begin
2   model2 = model |> cpu
3   if hasproperty(tokenizer, :cache)
4     # empty cache
5     tokenizer2 = similar(tokenizer)
6   end
7   output_path = joinpath(directory2, "model.bson")
8   history_path = joinpath(directory2, "history.json")
9   BSON.bson(
10    output_path,
11    Dict(
12      :model=> model2,
13      :tokenizer=>tokenizer,
14      :indexer=>indexer
15    )
16  )
17  println("saved model to $(output_path).")
18
19  open(history_path,"w") do f
20    JSON.print(f, history)
21  end
22  println("saved history to $(history_path).")
23
24 end
```

## Tip

Take note of the timestamped sub-directory so that you can load the saved model and parameters for use in the evaluation notebook.